

# Decision Tree | Assignment

**Question 1: What is a Decision Tree, and how does it work in the context of classification?**

**Ans:**

A Decision Tree is a supervised machine learning algorithm used for both classification and regression problems. It represents decisions in the form of a tree-like structure, where internal nodes represent conditions on features, branches represent decision rules, and leaf nodes represent the final output or class label.

## How a Decision Tree Works in Classification-

In a classification problem, a Decision Tree works by repeatedly splitting the dataset into smaller subsets based on feature values.

## Step-by-step working:

- The algorithm starts at the root node containing the full dataset.
- It selects the best feature to split the data based on an impurity measure such as Gini Impurity or Entropy.
- The dataset is divided into subsets according to the selected feature.
- This process is repeated recursively for each subset.
- The splitting stops when a stopping condition is met (pure node, max depth reached, or minimum samples).
- Each leaf node assigns a class label based on the majority class of the samples in that node.

## Key Characteristics:

- Easy to understand and interpret
- Mimics human decision-making
- Works with both numerical and categorical data

**Question 2: Explain the concepts of Gini Impurity and Entropy as impurity measures. How do they impact the splits in a Decision Tree?**

**Ans:**

**Gini Impurity:** Gini Impurity measures the probability that a randomly chosen sample would be incorrectly classified.

**Formula:**  $\text{Gini} = 1 - \sum(p_i^2)$   
Where  $p_i$  is the probability of class  $i$ .

## Interpretation:

- $\text{Gini} = 0 \rightarrow \text{Node is pure}$
- Lower Gini  $\rightarrow$  Better split

**Entropy:** Entropy measures the level of disorder or randomness in the dataset.

**Formula:**  $\text{Entropy} = -\sum(p_i \log_2 p_i)$

**Interpretation:**

- Entropy = 0 → Perfectly pure node
- Higher entropy → More mixed classes

**Impact on Decision Tree Splits:**

- Both metrics aim to reduce impurity
- The algorithm chooses the split that gives maximum impurity reduction
- Gini is computationally faster
- Entropy is more sensitive to class imbalance

**Question 3: What is the difference between Pre-Pruning and Post-Pruning in Decision Trees? Give one practical advantage of using each.**

**Ans:**

1. **Pre-Pruning (Early Stopping):** Pre-pruning stops the tree growth early by setting constraints.

**Common techniques:**

- Maximum depth
- Minimum samples per split
- Minimum impurity decrease

**Practical Advantage:**

- Reduces overfitting
- Faster training time

2. **Post-Pruning:** Post-pruning allows the tree to grow fully and then trims unnecessary branches.

**Common techniques:**

- Cost complexity pruning
- Reduced error pruning

**Practical Advantage:**

- Often produces better generalization
- More accurate than pre-pruning

**Question 4: What is Information Gain in Decision Trees, and why is it important for choosing the best split?**

**Ans:**

Information Gain measures the reduction in entropy after a dataset is split on a feature.

**Formula:** Information Gain = Entropy(parent) – Weighted Entropy(children)

### Importance of Information Gain

- Helps identify the most informative feature
- Ensures optimal splits
- Improves classification accuracy
- Reduces randomness at each node

**Question 5: What are some common real-world applications of Decision Trees, and what are their main advantages and limitations?**

**Dataset Info:**

- Iris Dataset for classification tasks (`sklearn.datasets.load_iris()` or provided CSV).
- Boston Housing Dataset for regression tasks (`sklearn.datasets.load_boston()` or provided CSV).

**Ans:**

### Common Applications

- Healthcare: Diagnosing diseases based on patient symptoms.
- Finance: Credit scoring and fraud detection.
- Marketing: Customer churn prediction (identifying which customers might leave).

### Advantages

- Easy to interpret
- Handles non-linear relationships
- Requires minimal data preprocessing

### Limitations

- Prone to overfitting
- Unstable with small data changes
- Lower accuracy compared to ensemble models

**Question 6: Write a Python program to:**

- Load the Iris Dataset
- Train a Decision Tree Classifier using the Gini criterion
- Print the model's accuracy and feature importances

**Ans:**

```
from sklearn.datasets import load_iris  
from sklearn.model_selection import train_test_split  
from sklearn.tree import DecisionTreeClassifier
```

```

from sklearn.metrics import accuracy_score

# Load dataset
iris = load_iris()
X, y = iris.data, iris.target

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train model
model = DecisionTreeClassifier(criterion='gini', random_state=42)
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Accuracy
print("Accuracy:", accuracy_score(y_test, y_pred))

# Feature Importances
print("Feature Importances:", model.feature_importances_)

Output:

- Accuracy ≈ 1.0
- Displays importance of each feature

```

### **Question 7: Write a Python program to:**

- **Load the Iris Dataset**
- **Train a Decision Tree Classifier with max\_depth=3 and compare its accuracy to a fully-grown tree.**

### **Ans:**

```

# Limited depth model
model_limited = DecisionTreeClassifier(max_depth=3, random_state=42)
model_limited.fit(X_train, y_train)

# Fully grown model
model_full = DecisionTreeClassifier(random_state=42)
model_full.fit(X_train, y_train)

print("Accuracy (max_depth=3):", accuracy_score(y_test, model_limited.predict(X_test)))
print("Accuracy (full tree):", accuracy_score(y_test, model_full.predict(X_test)))

```

### **Observation:**

- Fully grown tree may overfit
- Limited depth provides better generalization

### **Question 8: Write a Python program to:**

- **Load the Boston Housing Dataset**
- **Train a Decision Tree Regressor**
- **Print the Mean Squared Error (MSE) and feature importances**

**Ans:**

```
from sklearn.datasets import load_boston
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error

boston = load_boston()
X, y = boston.data, boston.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model = DecisionTreeRegressor(random_state=42)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print("MSE:", mean_squared_error(y_test, y_pred))
print("Feature Importances:", model.feature_importances_)
```

### **Question 9: Write a Python program to:**

- **Load the Iris Dataset**
- **Tune the Decision Tree's max\_depth and min\_samples\_split using GridSearchCV**
- **Print the best parameters and the resulting model accuracy**

**Ans:**

```
from sklearn.model_selection import GridSearchCV

param_grid = {
    'max_depth': [2, 3, 4, 5, None],
    'min_samples_split': [2, 5, 10]
}

grid = GridSearchCV(DecisionTreeClassifier(random_state=42), param_grid, cv=5)
grid.fit(X_train, y_train)

print("Best Parameters:", grid.best_params_)
print("Best Accuracy:", grid.best_score_)
```

### **Question 10: Imagine you're working as a data scientist for a healthcare company that wants to predict whether a patient has a certain disease. You have a large dataset with**

**mixed data types and some missing values. Explain the step-by-step process you would follow to:**

- Handle the missing values
- Encode the categorical features
- Train a Decision Tree model
- Tune its hyperparameters
- Evaluate its performance And describe what business value this model could provide in the real-world

**Ans:**

## 1. Handling Missing Values

Missing values are common in healthcare datasets due to incomplete medical records or skipped tests.

### Steps to Handle Missing Values

#### a) Identify Missing Data

- Check missing values using summary statistics or functions like `isnull()` in Python.
- Understand which features have a high percentage of missing values.

#### b) Numerical Features

- Replace missing values with:
  - Mean if data is normally distributed
  - Median if data contains outliers
- Example: Missing blood pressure or cholesterol values can be filled with the median.

#### c) Categorical Features

- Replace missing values with:
  - Mode (most frequent category)
  - A new category such as "Unknown"
- This ensures no data is lost.

### Why This Step Is Important

- Prevents data loss
- Ensures the Decision Tree can process all records
- Improves model accuracy and stability

## 2. Encoding the Categorical Features

Machine learning models require numerical input, so categorical variables must be encoded.

### Encoding Techniques Used

#### a) Label Encoding

- Used for ordinal data (e.g., disease severity: mild < moderate < severe).
- Converts categories into integer values.

#### b) One-Hot Encoding

- Used for nominal data (e.g., gender, blood group).
- Creates binary columns for each category.

### Importance of Encoding

- Converts real-world medical information into machine-readable format
- Avoids introducing incorrect order in non-ordinal features

## 3. Training a Decision Tree Model

After preprocessing, the dataset is ready for model training.

### Steps to Train the Model

- Split the dataset into training and testing sets (e.g., 80% training, 20% testing).

- Choose a Decision Tree Classifier.
- Select an impurity measure:
  - Gini Impurity for faster computation
  - Entropy for more informative splits
- Train the model using the training dataset.

### **Why Decision Trees Are Suitable in Healthcare**

- Easy to interpret by doctors and clinicians
- Decision rules can be explained clearly
- Handles non-linear relationships effectively

## **4. Hyperparameter Tuning**

A fully grown Decision Tree may overfit the data, so tuning is essential.

### **Important Hyperparameters to Tune**

- max\_depth – Controls tree depth
- min\_samples\_split – Minimum samples required to split a node
- min\_samples\_leaf – Minimum samples required in a leaf node
- criterion – Gini or Entropy

### **Tuning Method**

- Use GridSearchCV with cross-validation
- Select the parameter combination that gives the best performance

### **Benefits of Tuning**

- Reduces overfitting
- Improves generalization to new patients
- Increases model reliability

## **5. Evaluating Model Performance**

Evaluating the model ensures it performs well on unseen data.

### **Key Evaluation Metrics**

- Accuracy – Overall correctness of predictions
- Precision – How many predicted disease cases are actually correct
- Recall (Sensitivity) – Ability to identify patients with the disease
- F1-score – Balance between precision and recall
- Confusion Matrix – Shows correct and incorrect predictions

### **Importance in Healthcare**

- High recall is crucial to avoid missing actual disease cases
- Balanced evaluation prevents harmful misdiagnosis

## **6. Business Value in the Real-World Healthcare Setting**

### **Key Business and Clinical Benefits**

- **Early Disease Detection**
  - Helps doctors identify high-risk patients early
- **Improved Patient Outcomes**
  - Early treatment leads to better recovery rates
- **Cost Reduction**
  - Prevents expensive late-stage treatments
- **Decision Support for Doctors**
  - Provides explainable insights for clinical decisions
- **Scalability**
  - Can analyze thousands of patient records quickly